

When there is a large error in its predictions, you should try:

- ① Get more training examples.
- ② Try smaller sets of features
- ③ Try getting additional features
- ④ Try getting polynomial features (x_1^2, x_2^2, x_1x_2)
- ⑤ Try decreasing/increasing λ .

we shouldn't randomly choose them to implement - machine learning diagnostic

Evaluate a Hypothesis \rightarrow overfitting.

split training set into training set and test (validation) set, (70% to 30%).

- Learning parameter θ from training data (minimize $J(\theta)$)

- Compute test error.

$$\text{linear regression: } J_{\text{test}}(\theta) = \frac{1}{m_{\text{test}}} \sum_{i=1}^{m_{\text{test}}} (\text{hor}x_i^{(i)}) - y_i^{(i)})^2$$

$$\text{logistic regression (classification): } J_{\text{test}}(\theta) = -\frac{1}{m_{\text{test}}} \sum_{i=1}^{m_{\text{test}}} y_i^{(i)} \log(h_{\text{test}}(x_i^{(i)})) + (1-y_i^{(i)}) \log(1-h_{\text{test}}(x_i^{(i)}))$$

- Misclassification error (0/1 missclassification error).

$$\text{error}(h(x), y) = \begin{cases} 1 & \text{if } h(x) \geq 0.5 \text{ and } y=0 \text{ or if } h(x) \leq 0.5 \text{ and } y=1} \\ 0 & \text{otherwise.} \end{cases}$$

$$\text{Test error} = \frac{1}{m_{\text{test}}} \sum_{i=1}^{m_{\text{test}}} \text{error}(h(x_i^{(i)}), y_i^{(i)})$$

Model selection:

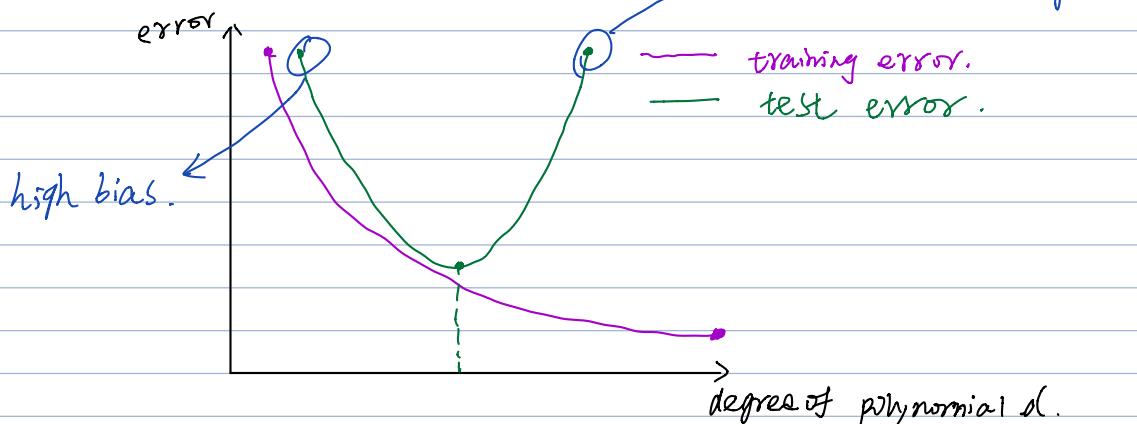
split training sets into training, validation, test sets.

- ① optimize the parameters in θ using the training set for each polynomial degree
- ② Find the polynomial degree d with least error using the cross validation set.
- ③ Estimate the generalization error using the test set with $J_{\text{test}}(\theta^{(d)})$

Bias and variance.

High bias \rightarrow underfit

High variance \rightarrow overfit.



Bias problem (underfit) :

$J_{\text{train}}(\theta)$ will be high
 $J_{\text{cv}}(\theta) \approx J_{\text{train}}(\theta)$

Varionce (overfit) :

Regularization & bias/variance.

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (\theta_0 x_i^{(1)} + \theta_1 x_i^{(2)})^2 + \frac{\lambda}{2m} \sum_{j=1}^m \theta_j^2$$

large $\lambda \rightarrow$ underfitting \rightarrow high bias.

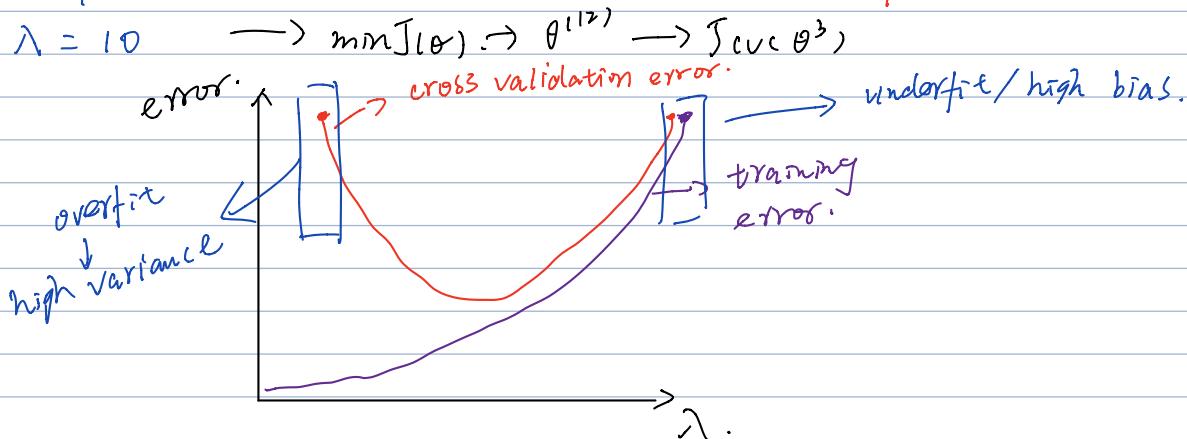
small $\lambda \rightarrow$ overfitting \rightarrow high variance.

define $I_{train}(\theta)$ without regularization

$$\text{J}_{\text{train}}(\theta) = \frac{1}{m} \sum_{i=1}^m (\hat{y}_i(x_i^{(i)}) - y_i^{(i)})^2 = \text{J}_{\text{CV}}(\theta) = \text{J}_{\text{test}}(\theta).$$

training error 不包含正则化.

\Rightarrow pick $\theta^{(k)}$ with lowest
 $J_{CV}(\theta^{(k)})$.

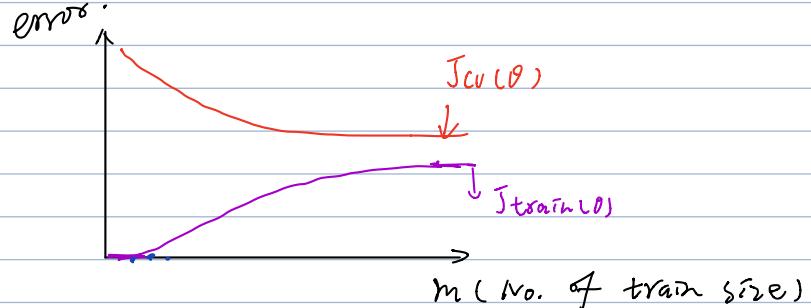


Learning curves.

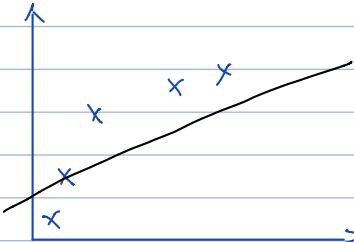
Suppose $h_0(x) = \theta_0 + \theta_1 x + \theta_2 x^2$

$$J_{\text{train}}(\theta) = \frac{1}{2m} \sum_{i=1}^m (\theta_0 x^{(i)}) - y^{(i)} \rightarrow$$

$$J_{CV}(\theta) = \frac{1}{2m} \sum_{i=1}^m (h\theta(x^{(i)}) - y^{(i)})^2$$



High bias: $h(x) = \theta_0 + \theta_1 x$. \rightarrow underfit -



desired performance.

error

$J_{CV}(\theta)$

$J_{train}(\theta)$

=> high $J_{CV}(\theta)$
& $J_{train}(\theta)$.

m (No. of train size)

If a learning algorithm is suffering from high bias, getting more training data will not help much.

High variance: $h(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \dots + \theta_{100} x^{100}$



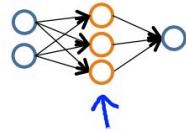
增加训练集可以提高高模型效果。

$J_{train}(\theta) < J_{CV}(\theta)$ but the difference between them remains significant.

- ① Get more training examples. \rightarrow fix high variance
- ② Try smaller sets of features \rightarrow fix high variance
- ③ Try getting additional features \rightarrow fix high bias.
- ④ Try getting polynomial features ($x_1^2, x_2^2, x_1 x_2$) \leftarrow fix high bias.
- ⑤ Try decreasing/increasing λ .
 ↓
 fix high bias fix high variance.

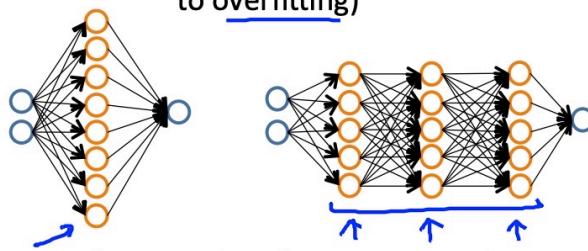
Neural networks and overfitting

→ “Small” neural network
(fewer parameters; more prone to underfitting)



Computationally cheaper

→ “Large” neural network
(more parameters; more prone to overfitting)



Computationally more expensive.

Use regularization (λ) to address overfitting.

$J_{reg}(\theta)$

通常选择较大的神经网络并使用正则化以使用较少的神经元效果好。

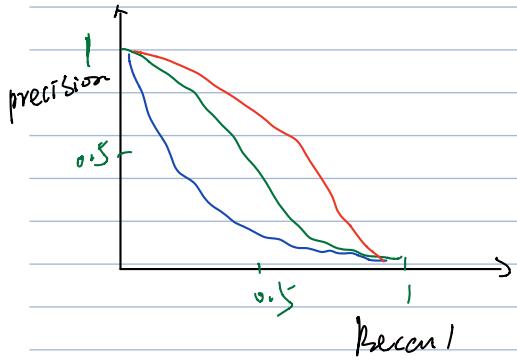
True Class			
		1	0
prediction class	1	True positive (TP) False positive.	False (FP)
	0	False negative (FN)	True (TN) negative

precision = $\frac{TP}{TP + FP}$ ↗ 在预测出有 tumor 的人中,
 真正有 tumor 的病人

Recall = $\frac{TP}{TP + FN}$ ↗ 在实际有 tumor 的人中,
 预测对的人数.

越高越好!

Trade off precision and recall.



F1 Score (F score)

$$2 \cdot \frac{PR}{P+R}$$